

FICTIONAL CASE · SYNTHETIC DATA

CENTERLINE BANK · CLAUDE PILOT

AI the RMs actually use, and the CCO trusts

Built, not described · Measured, not promised · Compliant by design

Technical-challenge presentation · **Commercial Banking**
Senior Forward Deployed Engineer candidate

Franck Benichou · June 2026

About me

SNAPSHOT



Franck Benichou

Senior AI Engineer · Carta (LPPA) · Toronto

- **AI Engineering — Document Intelligence**
- **Data Science + Consulting / start-up** — Intact · EY · Deloitte · Carta
- Technical reviewer · [DeepSeek in Practice](#) (Packt, 2025)
- Founder · **HEC Forecast** student data-science conference (2018 → present)
- Français · English · Español

WHY CAYLENT'S SENIOR FDE — AND WHY ME

- 1 **It's the work I already do.** AWS + Anthropic is my current stack — I build with the Anthropic APIs and Claude Code every day at Carta, in the document-intelligence team.
- 2 **Why I'm excited.** I'm an AI builder who wants direct commercial impact on the front lines — prototyping and taking things into production.
- 3 **Why I'm a fit.** Beyond building, my unconventional, hybrid background brings judgement and client-facing strength — not just code.

The situation

Centerline's Commercial Banking division bought Claude for Work — and deployed nothing **yet**.

\$350M

portfolio managed by **3 relationship managers**

30–40%

of RM time on **documentation, covenant tracking & call prep**

0

workflows in production — **licenses bought, not used**

The CCO, **Rachel Torres**, has one stated priority: an **early-warning** capability for contact gaps and financial deterioration — RM-private, compliance-gated. *That's my mandate for the next 90 days.*

Meet the three RMs

Three RMs, three very different relationships with AI. Adoption is the real risk — so we design for each, not for the org chart.

Sarah Okonkwo

THE LINCHPIN · NATURAL TEACHER

Book: Meridian · BlueLine — both in covenant breach

Diligent logger. Burned once by an AI that gave **wrong covenant info** — so trust is existential.

If she adopts, she trains the team.

Tom DeLuca

THE BUILDER · POWER-USER

Book: Crestwood (A-rated) · Arcadia (construction)

Built 5 AI workflows on personal ChatGPT. Wants to **co-design**, not be policed.

Already proved AI helps — by breaking policy.

Marcus Webb

THE SKEPTIC · 22-YR VETERAN

Book: Summit — DSCR drifting toward the floor

Mentors 4 juniors. **~71 annual memos/yr.** Logs sparsely; trusts his own memory.

Win him and the rollout tips.

The catalyst — one RM didn't wait

IT AUDIT FINDING

Tom built 5 AI workflows on personal ChatGPT

Client financials and names pasted into a personal account — **two compliance violations**. Draw-response letters · relationship memos · client emails · call-prep briefs · a weekly portfolio summary.

THE REFRAME

- **Right instinct, wrong tools.** Tom proved AI helps here — by breaking policy to do it.
- **Starting point, not blueprint.** We keep the instinct, fix the violations, and turn one person's hack into a governed, shareable library.

Compliance isn't a wrapper — it's the architecture

Section 4 is trickier than it looks. Five constraints decided every design choice that follows.

§4.2

No creditworthiness language — even implicitly. A disclaimer does **not** count.

§2.1

Credit grade / watchlist can't even be **inputs** to the AI.

§5

No AI on **Special-Assets** or litigation credits.

§4.1

Automated monitoring needs **CCO approval** to launch.

§4.3

Human review before anything is client-facing or filed.

The credit **judgment must originate with the RM**. Claude organizes facts and flags gaps — it never characterizes credit. And we enforce that **in code, not in a prompt**.

The signal's already there — the record can't be trusted

Eight files, current through **2025-05-31**. The hard part isn't finding data — the system of record is wrong three ways: **mis-dated logs** · **decisions that live only in email** · **status labels that lag the trend**.

Meridian

A careful action plan in the memo — that the data shows was **never executed**.

Crestwood

Attrition risk on a pristine, healthy borrower — invisible to every distress metric.

Across the book

Engagement **diverges from distress** — RMs go quiet on the credits that are sliding.

So the job isn't to find hidden data. It's to **organize scattered facts and make the record trustworthy** — and leave the judgment with the RM. *That's the spine of both demos.*